


Open Challenges in Making AI Autonomy Reliable

I. Long-Horizon & Rollback

Detect drift early. Recover safely.



Goal: Maintain progress over long horizons and recover cleanly from deviations.

II. Safety & Governance

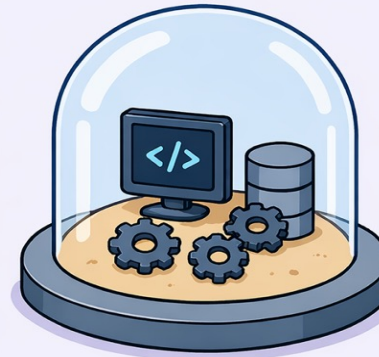
Operate within bounds. Prove it.

Permission Boundary



Restrict actions to authorized resources and capabilities.

Sandbox Environment



Isolate execution to limit blast radius of failures.

Audit Trail



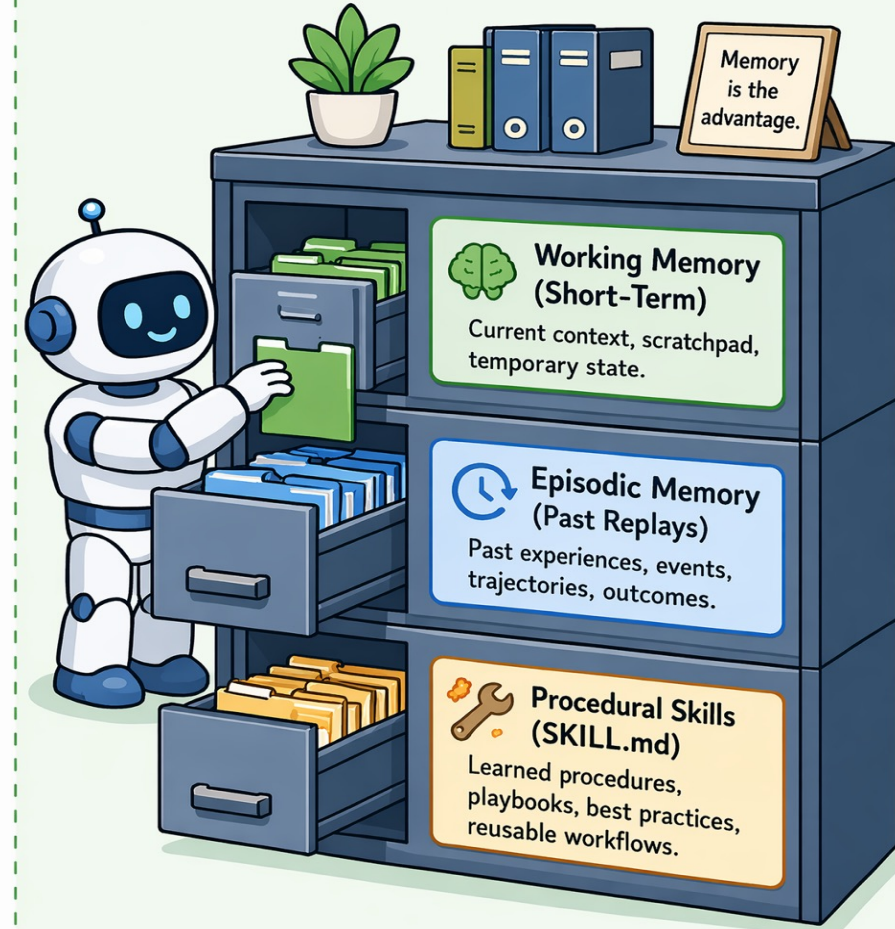
- 🕒 Who did what
- 📍 When & where
- ⚙️ Inputs & tools
- ✅ Outcomes
- 🛡️ Policy checks



Goal: Ensure operational safety, enforce policy, and enable accountability through evidence.

III. Persistent Memory

Remember across time. Use appropriately.



Goal: Manage different memory types across timescales to enable continuity and competence.